

Housing_Predictions

April 22, 2018

1 Predicting Housing Prices in Ames, Iowa

1.1 Introduction

When a person is thinking about what house to buy, what are the features that come to mind? Some people want to make sure that they have a basement. Some want to make sure that they have multiple rooms. Some want to make sure that they have multiple bathrooms. Of course, there is the real estate mantra that all that matters is "Location. Location. Location." The objective of this report is to explore what features drive the Sale Price of a house the most.

1.2 Initial Data Exploration

When doing analysis with python, the first step is to load the modules that will be used. Some of the main modules are pandas which is used to manipulate datasets, numpy which is useful for some mathematical computations, matplotlib which is used for graphing and sklearn which is used for predictive modeling after the data is preprocessed.

Of course, no data can be analyzed if there is no data present. So, the first step in analyzing the data is to load the data. The data for this report is split between a train set and a test set. The train and the test set are combined into a data set and stored in a variable called `houses_full`. The train and the test are combined so that when variables are transformed and engineered, they will be transformed in both the train sets and the test sets. The train set contains the data of 81 different features for 1460 houses. The test set contains the data for 1461 houses containing the same features. Some of the features for these houses include the square footage of the entire house, the square footage of each floor including the basement if there's a basement, the number of fireplaces, the neighborhood, and the size of the garage if the house has one. All attributes are available in the Appendix #1c-e.

Further exploration reveals that both the train and test sets have missing values. However, most of these missing value are for features such as pool and fence. The obvious reason that these values are missing for these features is that these houses don't have pools or fences. In Appendix #4 all missing values for categorical variables are replaced with the word "MISSING". Categorical variables are variables that are not represented by numbers. For example, the neighborhood where the house is located is a categorical variable. Unfortunately the test set contains missing numerical data for the Masonry Veneer Area variable. In Appendix #9b all missing numerical entries are replaced with the mean of the feature.

1.3 Transforming the Data

Unfortunately, in order to make accurate predictions on the test dataset, the data cannot be used in its original state. There are some necessary transformations that need to be made in order to give the most accurate predictions. So the first step in the data transformation is to transform skewed variables. A skew can be seen in the target variable SalePrice. In the graph in Appendix #3a the positive skew is evident. The positive skew means that most of the prices for the houses are lower in price while there is a smaller amount of home prices being more expensive. For analysis purposes in Appendix #3c, a log+1 transformation is applied to all numeric variables including the SalePrice variable. The result is that all numeric variables will be more normally distributed. This means that most home prices will have a middle value with the other prices being equally more expensive or less expensive. Essentially, the new numeric data is more evenly distributed. This effect can be seen on the SalePrice variable in Appendix #3d.

The next step in the transformation is to change the categorical variables. So all variables that are not represented as a numeric value, will be converted to a numeric value. For example, the name of the neighborhood where the house is located will no longer be represented by a name. It will be represented as a number. This transformation is shown in Appendix #5. Not only are these variables converted to numbers, but they're labels now include _E in order to distinguish them from the original variables. The effect of the transformation is shown in Appendix #5a-3.

Once these transformations are complete, the next step in building a predictive model is to select the correct features from the dataset. The dataset starts with 81 different variables. However, not all of these variables will be helpful when trying to predict the price of a home. So correlations between the SalePrice variable and the other variables are checked, in order to discover which features will help predict the value of a home. The correlations between SalePrice and the other variables are shown in Appendix #6b - Appendix #6g. The top correlated categorical variables are shown in Appendix #6b, #6d and #6e. The top numerical variables are shown in Appendix #6c, #6f and #6g. These correlation tables show that the top correlated features are the overall quality of the house, the square footage of the living area, the neighborhood, the amount of cars that fit in the garage and lots of other features. The top 25 categorical features and the top 25 numerical features are selected from the full dataset, train set and test set. These are the variables that will be used when building the predictive model.

It would be great if all the variables that are present in the dataset are the only variables that would help predict home prices. Unfortunately some data will need to be "engineered" or created in order to create the most accurate model. Some of the features that can be created are created from some features that are collinear. Collinear features are features that pretty evenly move with the target variable. For example, the square footage of the first floor is generally close to the square footage of the second floor if there is one present. So the square footage of the first and second floor are said to be collinear. However, having a second floor may change the price of a house. Having a basement may change the price of the house as well. Appendix #7b shows the creation of features to show whether the house, has a basement, has a garage, has a second floor, has an open porch, has a screen porch, has a fire place, has a masonry veneer, has a wood deck, has a pool, has a basement bathroom, has an extra bathroom, or is new.

With the data transformed, new data created, and the appropriate features selected, there is only one step left before the predictive models can be created. That is to remove outliers. Since, the square footage of the general living area is an important variable, it would be helpful to remove outliers. There are two obvious outliers, which can be seen in Appendix #8a. These outliers are removed and the data is ready for modeling.

1.4 Predictive Modeling

The first step in creating the predictive model is to get dummy variables for the categorical features. Explaining a dummy variable is best done with an example. For the neighborhood variable, for example, each neighborhood becomes its own column. If the house is in that neighborhood, the dummy variable will show a 1. If the house is not in that neighborhood, the dummy variable will show a 0. This step is shown in Appendix #9a. Once the dummy variables are created, the data is partitioned. The data that is going to be used to build the model is stored in the variable `train_subset`. The data stored in this variable is separated into its own train set and test set. The train set contains 70% of the data in `train_subset` and will be used to build the model. The test set contains the remaining 30% of the data and will be used to test the model.

Next four different machine learning algorithms are used to model the data. These models include a Linear Regression model, a Lasso model, a Ridge model and an ElasticNet model. Each of these models has a graph accompanying it, which compares the predictions with the actual selling price. The Lasso model is the final choice for making predictions on the competition test set.

1.5 Conclusion

After doing this thorough analysis on the housing data in Ames, Iowa, a few conclusions can be made. In the predictive modeling section of this report, the features that have the most effect on the price of a house can be determined. The graphs in Appendix #9 show which variables have the greatest impact. This analysis determined that the selling price of a home in Ames, Iowa is most influenced by the size of the general living area, whether the house has a basement, deck, second floor, extra bathrooms, a garage, a fire place or open porch, the size of additional floors and basements, and the overall quality of the house. Although the real estate agents are right when it comes to "Location. Location. Location," this analysis shows that there are more important criteria when determining the price of a house.

1.6 Appendix

1.6.1 1. Initial Data Exploration

1a. Loading Modules

```
In [119]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats
import seaborn as sns
from IPython.display import HTML, display
from sklearn.manifold import TSNE
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, RidgeCV, LassoCV, ElasticNetCV, L
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, make_scorer
```

```
from sklearn.preprocessing import StandardScaler
%matplotlib inline
```

1b. Loading Data

```
In [120]: houses_train = pd.read_csv('train.csv')
          houses_test = pd.read_csv('test.csv')
          houses_full = houses_train.append(houses_test, ignore_index = True)
          houses = houses_full[ :1460]
          test = houses_full[1460: ]
```

1c. Inspecting Data

1c-1. Training Data

```
In [121]: houses.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
1stFlrSF          1460 non-null int64
2ndFlrSF          1460 non-null int64
3SsnPorch         1460 non-null int64
Alley             91 non-null object
BedroomAbvGr      1460 non-null int64
BldgType          1460 non-null object
BsmtCond          1423 non-null object
BsmtExposure      1422 non-null object
BsmtFinSF1        1460 non-null float64
BsmtFinSF2        1460 non-null float64
BsmtFinType1      1423 non-null object
BsmtFinType2      1422 non-null object
BsmtFullBath      1460 non-null float64
BsmtHalfBath      1460 non-null float64
BsmtQual          1423 non-null object
BsmtUnfSF         1460 non-null float64
CentralAir        1460 non-null object
Condition1        1460 non-null object
Condition2        1460 non-null object
Electrical        1459 non-null object
EnclosedPorch     1460 non-null int64
ExterCond         1460 non-null object
ExterQual         1460 non-null object
Exterior1st       1460 non-null object
Exterior2nd       1460 non-null object
Fence             281 non-null object
FireplaceQu       770 non-null object
Fireplaces        1460 non-null int64
```

Foundation	1460	non-null	object
FullBath	1460	non-null	int64
Functional	1460	non-null	object
GarageArea	1460	non-null	float64
GarageCars	1460	non-null	float64
GarageCond	1379	non-null	object
GarageFinish	1379	non-null	object
GarageQual	1379	non-null	object
GarageType	1379	non-null	object
GarageYrBltd	1379	non-null	float64
GrLivArea	1460	non-null	int64
HalfBath	1460	non-null	int64
Heating	1460	non-null	object
HeatingQC	1460	non-null	object
HouseStyle	1460	non-null	object
Id	1460	non-null	int64
KitchenAbvGr	1460	non-null	int64
KitchenQual	1460	non-null	object
LandContour	1460	non-null	object
LandSlope	1460	non-null	object
LotArea	1460	non-null	int64
LotConfig	1460	non-null	object
LotFrontage	1201	non-null	float64
LotShape	1460	non-null	object
LowQualFinSF	1460	non-null	int64
MSSubClass	1460	non-null	int64
MSZoning	1460	non-null	object
MasVnrArea	1452	non-null	float64
MasVnrType	1452	non-null	object
MiscFeature	54	non-null	object
MiscVal	1460	non-null	int64
MoSold	1460	non-null	int64
Neighborhood	1460	non-null	object
OpenPorchSF	1460	non-null	int64
OverallCond	1460	non-null	int64
OverallQual	1460	non-null	int64
PavedDrive	1460	non-null	object
PoolArea	1460	non-null	int64
PoolQC	7	non-null	object
RoofMatl	1460	non-null	object
RoofStyle	1460	non-null	object
SaleCondition	1460	non-null	object
SalePrice	1460	non-null	float64
SaleType	1460	non-null	object
ScreenPorch	1460	non-null	int64
Street	1460	non-null	object
TotRmsAbvGrd	1460	non-null	int64
TotalBsmtSF	1460	non-null	float64

```

Utilities          1460 non-null object
WoodDeckSF         1460 non-null int64
YearBuilt          1460 non-null int64
YearRemodAdd       1460 non-null int64
YrSold             1460 non-null int64
dtypes: float64(12), int64(26), object(43)
memory usage: 924.0+ KB

```

1c-2. Test Data

```
In [122]: test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1459 entries, 1460 to 2918
Data columns (total 81 columns):
1stFlrSF          1459 non-null int64
2ndFlrSF          1459 non-null int64
3SsnPorch         1459 non-null int64
Alley             107 non-null object
BedroomAbvGr      1459 non-null int64
BldgType          1459 non-null object
BsmtCond          1414 non-null object
BsmtExposure      1415 non-null object
BsmtFinSF1        1458 non-null float64
BsmtFinSF2        1458 non-null float64
BsmtFinType1      1417 non-null object
BsmtFinType2      1417 non-null object
BsmtFullBath      1457 non-null float64
BsmtHalfBath      1457 non-null float64
BsmtQual          1415 non-null object
BsmtUnfSF         1458 non-null float64
CentralAir        1459 non-null object
Condition1        1459 non-null object
Condition2        1459 non-null object
Electrical        1459 non-null object
EnclosedPorch     1459 non-null int64
ExterCond         1459 non-null object
ExterQual         1459 non-null object
Exterior1st       1458 non-null object
Exterior2nd       1458 non-null object
Fence             290 non-null object
FireplaceQu       729 non-null object
Fireplaces        1459 non-null int64
Foundation        1459 non-null object
FullBath          1459 non-null int64
Functional        1457 non-null object
GarageArea        1458 non-null float64

```

GarageCars	1458	non-null	float64
GarageCond	1381	non-null	object
GarageFinish	1381	non-null	object
GarageQual	1381	non-null	object
GarageType	1383	non-null	object
GarageYrBltd	1381	non-null	float64
GrLivArea	1459	non-null	int64
HalfBath	1459	non-null	int64
Heating	1459	non-null	object
HeatingQC	1459	non-null	object
HouseStyle	1459	non-null	object
Id	1459	non-null	int64
KitchenAbvGr	1459	non-null	int64
KitchenQual	1458	non-null	object
LandContour	1459	non-null	object
LandSlope	1459	non-null	object
LotArea	1459	non-null	int64
LotConfig	1459	non-null	object
LotFrontage	1232	non-null	float64
LotShape	1459	non-null	object
LowQualFinSF	1459	non-null	int64
MSSubClass	1459	non-null	int64
MSZoning	1455	non-null	object
MasVnrArea	1444	non-null	float64
MasVnrType	1443	non-null	object
MiscFeature	51	non-null	object
MiscVal	1459	non-null	int64
MoSold	1459	non-null	int64
Neighborhood	1459	non-null	object
OpenPorchSF	1459	non-null	int64
OverallCond	1459	non-null	int64
OverallQual	1459	non-null	int64
PavedDrive	1459	non-null	object
PoolArea	1459	non-null	int64
PoolQC	3	non-null	object
RoofMatl	1459	non-null	object
RoofStyle	1459	non-null	object
SaleCondition	1459	non-null	object
SalePrice	0	non-null	float64
SaleType	1458	non-null	object
ScreenPorch	1459	non-null	int64
Street	1459	non-null	object
TotRmsAbvGrd	1459	non-null	int64
TotalBsmtSF	1458	non-null	float64
Utilities	1457	non-null	object
WoodDeckSF	1459	non-null	int64
YearBuilt	1459	non-null	int64
YearRemodAdd	1459	non-null	int64

```
YrSold          1459 non-null int64
dtypes: float64(12), int64(26), object(43)
memory usage: 923.4+ KB
```

1c-3. Full Data Set

```
In [123]: houses_full.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2919 entries, 0 to 2918
Data columns (total 81 columns):
1stFlrSF          2919 non-null int64
2ndFlrSF          2919 non-null int64
3SsnPorch         2919 non-null int64
Alley             198 non-null object
BedroomAbvGr      2919 non-null int64
BldgType          2919 non-null object
BsmtCond          2837 non-null object
BsmtExposure      2837 non-null object
BsmtFinSF1        2918 non-null float64
BsmtFinSF2        2918 non-null float64
BsmtFinType1      2840 non-null object
BsmtFinType2      2839 non-null object
BsmtFullBath      2917 non-null float64
BsmtHalfBath      2917 non-null float64
BsmtQual          2838 non-null object
BsmtUnfSF         2918 non-null float64
CentralAir        2919 non-null object
Condition1        2919 non-null object
Condition2        2919 non-null object
Electrical        2918 non-null object
EnclosedPorch     2919 non-null int64
ExterCond         2919 non-null object
ExterQual         2919 non-null object
Exterior1st       2918 non-null object
Exterior2nd       2918 non-null object
Fence             571 non-null object
FireplaceQu       1499 non-null object
Fireplaces        2919 non-null int64
Foundation        2919 non-null object
FullBath          2919 non-null int64
Functional        2917 non-null object
GarageArea        2918 non-null float64
GarageCars        2918 non-null float64
GarageCond        2760 non-null object
GarageFinish      2760 non-null object
GarageQual        2760 non-null object
```


GarageType	2762	non-null	object
GarageYrBltd	2760	non-null	float64
GrLivArea	2919	non-null	int64
HalfBath	2919	non-null	int64
Heating	2919	non-null	object
HeatingQC	2919	non-null	object
HouseStyle	2919	non-null	object
Id	2919	non-null	int64
KitchenAbvGr	2919	non-null	int64
KitchenQual	2918	non-null	object
LandContour	2919	non-null	object
LandSlope	2919	non-null	object
LotArea	2919	non-null	int64
LotConfig	2919	non-null	object
LotFrontage	2433	non-null	float64
LotShape	2919	non-null	object
LowQualFinSF	2919	non-null	int64
MSSubClass	2919	non-null	int64
MSZoning	2915	non-null	object
MasVnrArea	2896	non-null	float64
MasVnrType	2895	non-null	object
MiscFeature	105	non-null	object
MiscVal	2919	non-null	int64
MoSold	2919	non-null	int64
Neighborhood	2919	non-null	object
OpenPorchSF	2919	non-null	int64
OverallCond	2919	non-null	int64
OverallQual	2919	non-null	int64
PavedDrive	2919	non-null	object
PoolArea	2919	non-null	int64
PoolQC	10	non-null	object
RoofMatl	2919	non-null	object
RoofStyle	2919	non-null	object
SaleCondition	2919	non-null	object
SalePrice	1460	non-null	float64
SaleType	2918	non-null	object
ScreenPorch	2919	non-null	int64
Street	2919	non-null	object
TotRmsAbvGrd	2919	non-null	int64
TotalBsmtSF	2918	non-null	float64
Utilities	2917	non-null	object
WoodDeckSF	2919	non-null	int64
YearBuilt	2919	non-null	int64
YearRemodAdd	2919	non-null	int64
YrSold	2919	non-null	int64

dtypes: float64(12), int64(26), object(43)

memory usage: 1.8+ MB

1d. Viewing Numerical Data

1d-1. Training Data

```
In [124]: houses.describe()
```

```
Out [124]:
```

	1stFlrSF	2ndFlrSF	3SsnPorch	BedroomAbvGr	BsmtFinSF1	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	1162.626712	346.992466	3.409589	2.866438	443.639726	
std	386.587738	436.528436	29.317331	0.815778	456.098091	
min	334.000000	0.000000	0.000000	0.000000	0.000000	
25%	882.000000	0.000000	0.000000	2.000000	0.000000	
50%	1087.000000	0.000000	0.000000	3.000000	383.500000	
75%	1391.250000	728.000000	0.000000	3.000000	712.250000	
max	4692.000000	2065.000000	508.000000	8.000000	5644.000000	

	BsmtFinSF2	BsmtFullBath	BsmtHalfBath	BsmtUnfSF	EnclosedPorch	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	46.549315	0.425342	0.057534	567.240411	21.954110	
std	161.319273	0.518911	0.238753	441.866955	61.119149	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	223.000000	0.000000	
50%	0.000000	0.000000	0.000000	477.500000	0.000000	
75%	0.000000	1.000000	0.000000	808.000000	0.000000	
max	1474.000000	3.000000	2.000000	2336.000000	552.000000	

	...	OverallQual	PoolArea	SalePrice	ScreenPorch	\
count	...	1460.000000	1460.000000	1460.000000	1460.000000	
mean	...	6.099315	2.758904	180921.195890	15.060959	
std	...	1.382997	40.177307	79442.502883	55.757415	
min	...	1.000000	0.000000	34900.000000	0.000000	
25%	...	5.000000	0.000000	129975.000000	0.000000	
50%	...	6.000000	0.000000	163000.000000	0.000000	
75%	...	7.000000	0.000000	214000.000000	0.000000	
max	...	10.000000	738.000000	755000.000000	480.000000	

	TotRmsAbvGrd	TotalBsmtSF	WoodDeckSF	YearBuilt	YearRemodAdd	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	6.517808	1057.429452	94.244521	1971.267808	1984.865753	
std	1.625393	438.705324	125.338794	30.202904	20.645407	
min	2.000000	0.000000	0.000000	1872.000000	1950.000000	
25%	5.000000	795.750000	0.000000	1954.000000	1967.000000	
50%	6.000000	991.500000	0.000000	1973.000000	1994.000000	
75%	7.000000	1298.250000	168.000000	2000.000000	2004.000000	
max	14.000000	6110.000000	857.000000	2010.000000	2010.000000	

	YrSold
count	1460.000000

```

mean    2007.815753
std      1.328095
min     2006.000000
25%     2007.000000
50%     2008.000000
75%     2009.000000
max     2010.000000

```

[8 rows x 38 columns]

1d-2. Test Data

In [125]: test.describe()

```

Out [125]:
      1stFlrSF  2ndFlrSF  3SsnPorch  BedroomAbvGr  BsmtFinSF1  \
count  1459.000000  1459.000000  1459.000000  1459.000000  1458.000000
mean   1156.534613   325.967786    1.794380    2.854010   439.203704
std    398.165820   420.610226   20.207842    0.829788   455.268042
min     407.000000    0.000000    0.000000    0.000000    0.000000
25%     873.500000    0.000000    0.000000    2.000000    0.000000
50%    1079.000000    0.000000    0.000000    3.000000   350.500000
75%    1382.500000   676.000000    0.000000    3.000000   753.500000
max    5095.000000  1862.000000   360.000000    6.000000  4010.000000

      BsmtFinSF2  BsmtFullBath  BsmtHalfBath  BsmtUnfSF  EnclosedPorch  \
count  1458.000000  1457.000000  1457.000000  1458.000000  1459.000000
mean     52.619342    0.434454    0.065202   554.294925   24.243317
std    176.753926    0.530648    0.252468   437.260486   67.227765
min      0.000000    0.000000    0.000000    0.000000    0.000000
25%      0.000000    0.000000    0.000000   219.250000    0.000000
50%      0.000000    0.000000    0.000000   460.000000    0.000000
75%      0.000000    1.000000    0.000000   797.750000    0.000000
max    1526.000000    3.000000    2.000000  2140.000000  1012.000000

      ...  OverallQual  PoolArea  SalePrice  ScreenPorch  \
count  ...  1459.000000  1459.000000    0.0  1459.000000
mean   ...    6.078821    1.744345   NaN   17.064428
std    ...    1.436812   30.491646   NaN   56.609763
min    ...    1.000000    0.000000   NaN    0.000000
25%    ...    5.000000    0.000000   NaN    0.000000
50%    ...    6.000000    0.000000   NaN    0.000000
75%    ...    7.000000    0.000000   NaN    0.000000
max    ...   10.000000   800.000000   NaN   576.000000

      TotRmsAbvGrd  TotalBsmtSF  WoodDeckSF  YearBuilt  YearRemodAdd  \
count  1459.000000  1458.000000  1459.000000  1459.000000  1459.000000
mean     6.385195  1046.117970   93.174777  1971.357779  1983.662783
std     1.508895   442.898624  127.744882   30.390071   21.130467

```

min	3.000000	0.000000	0.000000	1879.000000	1950.000000
25%	5.000000	784.000000	0.000000	1953.000000	1963.000000
50%	6.000000	988.000000	0.000000	1973.000000	1992.000000
75%	7.000000	1305.000000	168.000000	2001.000000	2004.000000
max	15.000000	5095.000000	1424.000000	2010.000000	2010.000000

```

YrSold
count    1459.000000
mean     2007.769705
std       1.301740
min       2006.000000
25%       2007.000000
50%       2008.000000
75%       2009.000000
max       2010.000000

```

[8 rows x 38 columns]

1d-3. Full Data Set

In [126]: houses_full.describe()

```

Out[126]:
      1stFlrSF    2ndFlrSF    3SsnPorch  BedroomAbvGr  BsmtFinSF1  \
count  2919.000000  2919.000000  2919.000000   2919.000000  2918.000000
mean   1159.581706   336.483727    2.602261    2.860226   441.423235
std     392.362079   428.701456   25.188169    0.822693   455.610826
min     334.000000    0.000000    0.000000    0.000000    0.000000
25%     876.000000    0.000000    0.000000    2.000000    0.000000
50%    1082.000000    0.000000    0.000000    3.000000   368.500000
75%    1387.500000   704.000000    0.000000    3.000000   733.000000
max     5095.000000  2065.000000   508.000000    8.000000  5644.000000

      BsmtFinSF2  BsmtFullBath  BsmtHalfBath  BsmtUnfSF  EnclosedPorch  \
count  2918.000000  2917.000000  2917.000000  2918.000000  2919.000000
mean     49.582248    0.429894    0.061364   560.772104    23.098321
std    169.205611    0.524736    0.245687   439.543659    64.244246
min      0.000000    0.000000    0.000000    0.000000    0.000000
25%      0.000000    0.000000    0.000000   220.000000    0.000000
50%      0.000000    0.000000    0.000000   467.000000    0.000000
75%      0.000000    1.000000    0.000000   805.500000    0.000000
max    1526.000000    3.000000    2.000000  2336.000000   1012.000000

      ...  OverallQual  PoolArea  SalePrice  ScreenPorch  \
count  ...   2919.000000  2919.000000   1460.000000  2919.000000
mean   ...    6.089072    2.251799  180921.195890    16.062350
std    ...    1.409947   35.663946   79442.502883    56.184365
min    ...    1.000000    0.000000   34900.000000    0.000000
25%    ...    5.000000    0.000000  129975.000000    0.000000

```

50%	...	6.000000	0.000000	163000.000000	0.000000
75%	...	7.000000	0.000000	214000.000000	0.000000
max	...	10.000000	800.000000	755000.000000	576.000000

	TotRmsAbvGrd	TotalBsmtSF	WoodDeckSF	YearBuilt	YearRemodAdd	\
count	2919.000000	2918.000000	2919.000000	2919.000000	2919.000000	
mean	6.451524	1051.777587	93.709832	1971.312778	1984.264474	
std	1.569379	440.766258	126.526589	30.291442	20.894344	
min	2.000000	0.000000	0.000000	1872.000000	1950.000000	
25%	5.000000	793.000000	0.000000	1953.500000	1965.000000	
50%	6.000000	989.500000	0.000000	1973.000000	1993.000000	
75%	7.000000	1302.000000	168.000000	2001.000000	2004.000000	
max	15.000000	6110.000000	1424.000000	2010.000000	2010.000000	

	YrSold
count	2919.000000
mean	2007.792737
std	1.314964
min	2006.000000
25%	2007.000000
50%	2008.000000
75%	2009.000000
max	2010.000000

[8 rows x 38 columns]

1e. Top Rows of Data

In [127]: houses.head(10)

```
Out[127]:
```

	1stFlrSF	2ndFlrSF	3SsnPorch	Alley	BedroomAbvGr	BldgType	BsmtCond	\
0	856	854	0	NaN	3	1Fam	TA	
1	1262	0	0	NaN	3	1Fam	TA	
2	920	866	0	NaN	3	1Fam	TA	
3	961	756	0	NaN	3	1Fam	Gd	
4	1145	1053	0	NaN	4	1Fam	TA	
5	796	566	320	NaN	1	1Fam	TA	
6	1694	0	0	NaN	3	1Fam	TA	
7	1107	983	0	NaN	3	1Fam	TA	
8	1022	752	0	NaN	2	1Fam	TA	
9	1077	0	0	NaN	2	2fmCon	TA	

	BsmtExposure	BsmtFinSF1	BsmtFinSF2	...	SaleType	ScreenPorch	Street	\
0	No	706.0	0.0	...	WD	0	Pave	
1	Gd	978.0	0.0	...	WD	0	Pave	
2	Mn	486.0	0.0	...	WD	0	Pave	
3	No	216.0	0.0	...	WD	0	Pave	
4	Av	655.0	0.0	...	WD	0	Pave	

5	No	732.0	0.0	...	WD	0	Pave
6	Av	1369.0	0.0	...	WD	0	Pave
7	Mn	859.0	32.0	...	WD	0	Pave
8	No	0.0	0.0	...	WD	0	Pave
9	No	851.0	0.0	...	WD	0	Pave

	TotRmsAbvGrd	TotalBsmtSF	Utilities	WoodDeckSF	YearBuilt	YearRemodAdd	\
0	8	856.0	AllPub	0	2003	2003	
1	6	1262.0	AllPub	298	1976	1976	
2	6	920.0	AllPub	0	2001	2002	
3	7	756.0	AllPub	0	1915	1970	
4	9	1145.0	AllPub	192	2000	2000	
5	5	796.0	AllPub	40	1993	1995	
6	7	1686.0	AllPub	255	2004	2005	
7	7	1107.0	AllPub	235	1973	1973	
8	8	952.0	AllPub	90	1931	1950	
9	5	991.0	AllPub	0	1939	1950	

	YrSold
0	2008
1	2007
2	2008
3	2006
4	2008
5	2009
6	2007
7	2009
8	2008
9	2008

[10 rows x 81 columns]

1.6.2 2. Finding Missing Values

2a. Train Set

```
In [128]: train_nan = houses.isnull().sum().sort_values(ascending=False)
percent_nan = (houses.isnull().sum()/houses.isnull().count()).sort_values(ascending=False)
missing_table = pd.concat([train_nan, percent_nan], axis=1, keys=['Total', 'Percent'])
missing_table.head(25)
```

```
Out[128]:
```

	Total	Percent
PoolQC	1453	0.995205
MiscFeature	1406	0.963014
Alley	1369	0.937671
Fence	1179	0.807534
FireplaceQu	690	0.472603
LotFrontage	259	0.177397
GarageFinish	81	0.055479

GarageCond	81	0.055479
GarageYrBltd	81	0.055479
GarageQual	81	0.055479
GarageType	81	0.055479
BsmtExposure	38	0.026027
BsmtFinType2	38	0.026027
BsmtCond	37	0.025342
BsmtFinType1	37	0.025342
BsmtQual	37	0.025342
MasVnrArea	8	0.005479
MasVnrType	8	0.005479
Electrical	1	0.000685
Condition1	0	0.000000
Condition2	0	0.000000
2ndFlrSF	0	0.000000
3SsnPorch	0	0.000000
BedroomAbvGr	0	0.000000
GarageCars	0	0.000000

2b. Test Set

```
In [129]: test_nan = test.isnull().sum().sort_values(ascending=False)
percent_nan = (test.isnull().sum()/test.isnull().count()).sort_values(ascending=False)
missing_table = pd.concat([test_nan, percent_nan], axis=1, keys=['Total', 'Percent'])
missing_table.head(25)
```

```
Out [129]:
```

	Total	Percent
SalePrice	1459	1.000000
PoolQC	1456	0.997944
MiscFeature	1408	0.965045
Alley	1352	0.926662
Fence	1169	0.801234
FireplaceQu	730	0.500343
LotFrontage	227	0.155586
GarageYrBltd	78	0.053461
GarageQual	78	0.053461
GarageFinish	78	0.053461
GarageCond	78	0.053461
GarageType	76	0.052090
BsmtCond	45	0.030843
BsmtQual	44	0.030158
BsmtExposure	44	0.030158
BsmtFinType2	42	0.028787
BsmtFinType1	42	0.028787
MasVnrType	16	0.010966
MasVnrArea	15	0.010281
MSZoning	4	0.002742
BsmtFullBath	2	0.001371

BsmtHalfBath	2	0.001371
Functional	2	0.001371
Utilities	2	0.001371
Exterior2nd	1	0.000685

1.6.3 3. Transforming Skewed Data

3a. Sale Price Distribution

```
In [130]: ax = sns.distplot(houses['SalePrice'])
          ax.set_xlabel('Sale Price')
          ax.set_ylabel('Count')
          ax.set_title('Distribution of Sale Price')
          plt.show()
```



3b. Separating Variables

```
In [131]: categorical = houses_full.select_dtypes(include = ['object']).columns
          numerical = houses_full.select_dtypes(exclude = ["object"]).columns
          numerical = numerical.drop('Id')
          print('Categorical: ' + str(len(categorical)))
          print('Numerical: ' + str(len(numerical)))
```

Categorical: 43


Numerical: 37

3c. Transforming Skewed Features

```
In [132]: skewed = houses_full[numerical].apply(lambda x: stats.skew(x.dropna()))
          skewed = skewed[skewed > 0.75]
          skewed = skewed.index
          houses_full[skewed] = np.log1p(houses_full[skewed])
```

3d. Sale Price Distribution Transformed

```
In [133]: houses = houses_full[:1460]
          ax = sns.distplot(houses['SalePrice'])
          ax.set_xlabel('Sale Price Transformed')
          ax.set_ylabel('Count')
          ax.set_title('Distribution of Sale Price Transformed')
          plt.show()
```



1.6.4 4. Filling Missing Categorical Variables

```
In [134]: for feature in houses_full[categorical]:
          houses_full[feature] = houses_full[feature].astype('category')
          if houses_full[feature].isnull().any():
              houses_full[feature] = houses_full[feature].cat.add_categories(['MISSING'])
              houses_full[feature] = houses_full[feature].fillna('MISSING')

          houses_full.head()
```

```

Out[134]:   1stFlrSF  2ndFlrSF  3SsnPorch  Alley  BedroomAbvGr  BldgType  BsmtCond  \
0  6.753438  6.751101         0.0  MISSING              3    1Fam      TA
1  7.141245  0.000000         0.0  MISSING              3    1Fam      TA
2  6.825460  6.765039         0.0  MISSING              3    1Fam      TA
3  6.869014  6.629363         0.0  MISSING              3    1Fam      Gd
4  7.044033  6.960348         0.0  MISSING              4    1Fam      TA

   BsmtExposure  BsmtFinSF1  BsmtFinSF2  ...  SaleType  ScreenPorch  Street  \
0           No    6.561031         0.0  ...      WD          0.0    Pave
1           Gd    6.886532         0.0  ...      WD          0.0    Pave
2           Mn    6.188264         0.0  ...      WD          0.0    Pave
3           No    5.379897         0.0  ...      WD          0.0    Pave
4           Av    6.486161         0.0  ...      WD          0.0    Pave

   TotRmsAbvGrd  TotalBsmtSF  Utilities  WoodDeckSF  YearBuilt  YearRemodAdd  \
0      2.197225    6.753438    AllPub    0.000000    2003      2003
1      1.945910    7.141245    AllPub    5.700444    1976      1976
2      1.945910    6.825460    AllPub    0.000000    2001      2002
3      2.079442    6.629363    AllPub    0.000000    1915      1970
4      2.302585    7.044033    AllPub    5.262690    2000      2000

   YrSold
0   2008
1   2007
2   2008
3   2006
4   2008

[5 rows x 81 columns]

```

1.6.5 5. Encoding Categorical Variables

```

In [135]: def encode(frame, feature):
            ordering = pd.DataFrame()
            ordering['val'] = frame[feature].unique()
            ordering.index = ordering.val
            ordering['spmean'] = frame[[feature, 'SalePrice']].groupby(feature).mean()['SalePrice']
            ordering = ordering.sort_values('spmean')
            ordering['ordering'] = range(1, ordering.shape[0]+1)
            ordering = ordering['ordering'].to_dict()

            for cat, o in ordering.items():
                frame.loc[frame[feature] == cat, feature+'_E'] = o

qual_encoded = []
for q in categorical:
    encode(houses_full, q)
    qual_encoded.append(q+'_E')

```

```
print(qual_encoded)
```

```
['Alley_E', 'BldgType_E', 'BsmtCond_E', 'BsmtExposure_E', 'BsmtFinType1_E', 'BsmtFinType2_E',
```

5a. Viewing Encoding Transformation

5a-1. Train Set

```
In [136]: houses = houses_full[:1460]
          houses.describe()
```

```
Out[136]:
```

	1stFlrSF	2ndFlrSF	3SsnPorch	BedroomAbvGr	BsmtFinSF1	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	7.008452	2.864586	0.085679	2.866438	4.229731	
std	0.317431	3.293311	0.666876	0.815778	2.992052	
min	5.814131	0.000000	0.000000	0.000000	0.000000	
25%	6.783325	0.000000	0.000000	2.000000	0.000000	
50%	6.992096	0.000000	0.000000	3.000000	5.951943	
75%	7.238676	6.591674	0.000000	3.000000	6.569832	
max	8.453827	7.633370	6.232448	8.000000	8.638525	

	BsmtFinSF2	BsmtFullBath	BsmtHalfBath	BsmtUnfSF	EnclosedPorch	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	0.655398	0.425342	0.039486	5.648378	0.698019	
std	1.845045	0.518911	0.162599	1.854020	1.727317	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	5.411646	0.000000	
50%	0.000000	0.000000	0.000000	6.170651	0.000000	
75%	0.000000	1.000000	0.000000	6.695799	0.000000	
max	7.296413	3.000000	1.098612	7.756623	6.315358	

	...	MiscFeature_E	Neighborhood_E	PavedDrive_E	PoolQC_E	\
count	...	1460.000000	1460.000000	1460.000000	1460.000000	
mean	...	2.966438	12.860274	2.856164	1.008904	
std	...	0.205069	6.409677	0.496592	0.140703	
min	...	1.000000	1.000000	1.000000	1.000000	
25%	...	3.000000	7.000000	3.000000	1.000000	
50%	...	3.000000	13.000000	3.000000	1.000000	
75%	...	3.000000	17.000000	3.000000	1.000000	
max	...	5.000000	25.000000	3.000000	4.000000	

	RoofMatl_E	RoofStyle_E	SaleCondition_E	SaleType_E	Street_E	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	3.040411	2.608219	4.831507	5.204795	1.995890	
std	0.395845	1.209938	0.887253	0.929955	0.063996	
min	1.000000	1.000000	1.000000	1.000000	1.000000	
25%	3.000000	2.000000	5.000000	5.000000	2.000000	

50%	3.000000	2.000000	5.000000	5.000000	2.000000
75%	3.000000	2.000000	5.000000	5.000000	2.000000
max	8.000000	6.000000	6.000000	9.000000	2.000000

```

Utilities_E
count    1460.000000
mean      1.999315
std       0.026171
min       1.000000
25%      2.000000
50%      2.000000
75%      2.000000
max       2.000000

```

[8 rows x 81 columns]

5a-2. Test Set

```
In [137]: test = houses_full[1460:]
test.describe()
```

```

Out[137]:
```

	1stFlrSF	2ndFlrSF	3SsnPorch	BedroomAbvGr	BsmtFinSF1	\
count	1459.000000	1459.000000	1459.000000	1459.000000	1458.000000	
mean	6.999917	2.794559	0.046702	2.854010	4.223071	
std	0.327842	3.261396	0.493994	0.829788	2.971567	
min	6.011267	0.000000	0.000000	0.000000	0.000000	
25%	6.773652	0.000000	0.000000	2.000000	0.000000	
50%	6.984716	0.000000	0.000000	3.000000	5.862209	
75%	7.232372	6.517671	0.000000	3.000000	6.626049	
max	8.536211	7.529943	5.888878	6.000000	8.296796	

	BsmtFinSF2	BsmtFullBath	BsmtHalfBath	BsmtUnfSF	EnclosedPorch	\
count	1458.000000	1457.000000	1457.000000	1458.000000	1459.000000	
mean	0.707051	0.434454	0.044800	5.605885	0.816893	
std	1.912309	0.530648	0.172272	1.879851	1.814052	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	5.394761	0.000000	
50%	0.000000	0.000000	0.000000	6.133398	0.000000	
75%	0.000000	1.000000	0.000000	6.683048	0.000000	
max	7.331060	3.000000	1.098612	7.669028	6.920672	

	...	MiscFeature_E	Neighborhood_E	PavedDrive_E	PoolQC_E	\
count	...	1459.000000	1459.000000	1459.000000	1459.000000	
mean	...	2.967786	12.572995	2.805346	1.004798	
std	...	0.195080	6.532969	0.574204	0.114055	
min	...	1.000000	1.000000	1.000000	1.000000	
25%	...	3.000000	7.000000	3.000000	1.000000	
50%	...	3.000000	12.000000	3.000000	1.000000	

75%	...	3.000000	17.000000	3.000000	1.000000
max	...	4.000000	25.000000	3.000000	4.000000

	RoofMatl_E	RoofStyle_E	SaleCondition_E	SaleType_E	Street_E \
count	1459.000000	1459.000000	1459.000000	1459.000000	1459.000000
mean	3.019877	2.557916	4.833448	5.195339	1.995888
std	0.222875	1.177596	0.878226	0.947645	0.064018
min	3.000000	1.000000	1.000000	1.000000	1.000000
25%	3.000000	2.000000	5.000000	5.000000	2.000000
50%	3.000000	2.000000	5.000000	5.000000	2.000000
75%	3.000000	2.000000	5.000000	5.000000	2.000000
max	8.000000	6.000000	6.000000	10.000000	2.000000

	Utilities_E
count	1459.000000
mean	2.001371
std	0.037012
min	2.000000
25%	2.000000
50%	2.000000
75%	2.000000
max	3.000000

[8 rows x 81 columns]

5a-3. Full Data

In [138]: `houses_full.describe()`

```
Out[138]:
```

	1stFlrSF	2ndFlrSF	3SsnPorch	BedroomAbvGr	BsmtFinSF1 \
count	2919.000000	2919.000000	2919.000000	2919.000000	2918.000000
mean	7.004186	2.829584	0.066197	2.860226	4.226403
std	0.322650	3.277023	0.587089	0.822693	2.981325
min	5.814131	0.000000	0.000000	0.000000	0.000000
25%	6.776507	0.000000	0.000000	2.000000	0.000000
50%	6.987490	0.000000	0.000000	3.000000	5.912150
75%	7.235979	6.558198	0.000000	3.000000	6.598509
max	8.536211	7.633370	6.232448	8.000000	8.638525

	BsmtFinSF2	BsmtFullBath	BsmtHalfBath	BsmtUnfSF	EnclosedPorch \
count	2918.000000	2917.000000	2917.000000	2918.000000	2919.000000
mean	0.681207	0.429894	0.042140	5.627146	0.757435
std	1.878810	0.524736	0.167493	1.866773	1.771894
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	5.398163	0.000000
50%	0.000000	0.000000	0.000000	6.148468	0.000000
75%	0.000000	1.000000	0.000000	6.692703	0.000000
max	7.331060	3.000000	1.098612	7.756623	6.920672

	...	MiscFeature_E	Neighborhood_E	PavedDrive_E	PoolQC_E	\
count	...	2919.000000	2919.000000	2919.000000	2919.000000	
mean	...	2.967112	12.716684	2.830764	1.006852	
std	...	0.200105	6.472081	0.537299	0.128073	
min	...	1.000000	1.000000	1.000000	1.000000	
25%	...	3.000000	7.000000	3.000000	1.000000	
50%	...	3.000000	12.000000	3.000000	1.000000	
75%	...	3.000000	17.000000	3.000000	1.000000	
max	...	5.000000	25.000000	3.000000	4.000000	

	RoofMatl_E	RoofStyle_E	SaleCondition_E	SaleType_E	Street_E	\
count	2919.000000	2919.000000	2919.000000	2919.000000	2919.000000	
mean	3.030147	2.583076	4.832477	5.200069	1.995889	
std	0.321359	1.193942	0.882602	0.938689	0.063996	
min	1.000000	1.000000	1.000000	1.000000	1.000000	
25%	3.000000	2.000000	5.000000	5.000000	2.000000	
50%	3.000000	2.000000	5.000000	5.000000	2.000000	
75%	3.000000	2.000000	5.000000	5.000000	2.000000	
max	8.000000	6.000000	6.000000	10.000000	2.000000	


	Utilities_E
count	2919.000000
mean	2.000343
std	0.032062
min	1.000000
25%	2.000000
50%	2.000000
75%	2.000000
max	3.000000

[8 rows x 81 columns]

1.6.6 6. Selecting Features

6a. Boxplots of Categorical Variables


```
In [139]: def boxplot(x, y, **kwargs):
            sns.boxplot(x=x, y=y)
            x=plt.xticks(rotation=90)
            f = pd.melt(houses, id_vars=['SalePrice'], value_vars=categorical)
            g = sns.FacetGrid(f, col="variable", col_wrap=2, sharex=False, sharey=False, size=5)
            g = g.map(boxplot, "value", "SalePrice")
```



6b. Analysis of Variance (ANOVA)

```
In [140]: def anova(frame):
    anv = pd.DataFrame()
    anv['feature'] = categorical
    pvals = []
    for c in categorical:
        samples = []
        for cls in frame[c].unique():
            s = frame[frame[c] == cls]['SalePrice'].values
            samples.append(s)
        pval = stats.f_oneway(*samples)[1]
        pvals.append(pval)
    anv['pval'] = pvals
    return anv.sort_values('pval')

a = anova(houses)
a['disparity'] = np.log(1./a['pval'].values)
sns.barplot(data=a, x='feature', y='disparity')
plt.xticks(rotation=90)
```




6c. Spearman Correlation

```
In [141]: numerical_columns = houses.select_dtypes(include = ['int64', 'float64']).columns

def spearman(frame, features):
    spr = pd.DataFrame()
    spr['feature'] = features
    spr['spearman'] = [frame[f].corr(frame['SalePrice'], 'spearman') for f in features]
    spr = spr.sort_values('spearman')
    plt.figure(figsize=(6, 0.25*len(features)))
    sns.barplot(data=spr, y='feature', x='spearman', orient='h')


spearman(houses, numerical_columns)
```



6d. Sale Price Correlation Matrix (Categorical Variables)

```
In [142]: corr_matrix = houses[['SalePrice'] + qual_encoded].corr()
          f, ax = plt.subplots(figsize=(15,12))
          corr_graph = sns.heatmap(corr_matrix, vmax=0.8, square=True)
          corr_graph.set_title('Sale Price Correlation Matrix (Categorical Variables)')
```


```
Out[142]: Text(0.5,1,'Sale Price Correlation Matrix (Categorical Variables)')
```



6e. Sale Price Top 10 Correlation Matrix (Categorical Variables)

```
In [143]: k = 25
          cols = corr_matrix.nlargest(k, 'SalePrice')['SalePrice'].index
```


```
f, ax = plt.subplots(figsize=(30,25))
cm = np.corrcoef(houses[cols].values.T)
sns.set(font_scale=1.25)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_kws={'size':
        yticklabels=cols.values, xticklabels=cols.values})
hm.set_title('Sale Price Top Correlations (Categorical)')
plt.show()
```



6f. Sale Price Correlation Matrix (Numeric Variables)


```
In [144]: corr_matrix_2 = houses[numerical].corr()
f, ax = plt.subplots(figsize=(15,12))
corr_graph = sns.heatmap(corr_matrix_2, vmax=0.8, square=True)
corr_graph.set_title('Sale Price Correlation Matrix (Numeric Variables)')
```

```
Out[144]: Text(0.5,1,'Sale Price Correlation Matrix (Numeric Variables)')
```



6g. Sale Price Top 10 Correlation Matrix (Numeric Variables)

```
In [145]: k = 25
cols_2 = corr_matrix_2.nlargest(k, 'SalePrice')['SalePrice'].index
f, ax = plt.subplots(figsize=(30,25))
cm_2 = np.corrcoef(houses[cols_2].values.T)
sns.set(font_scale=1.25)
hm_2 = sns.heatmap(cm_2, cbar=True, annot=True, square=True, fmt='.2f', annot_kws={
    'yticklabels=cols_2.values, xticklabels=cols_2.values)
hm_2.set_title('Sale Price Top Correlations (Numeric)')
plt.show()
```



6h. Selecting Features Correlated with SalePrice

6h-1 Full Subset

```
In [146]: cat_subset = houses_full[cols]
cat_subset = cat_subset.drop('SalePrice', axis=1)
num_subset = houses_full[cols_2]
full_subset = pd.concat([num_subset, cat_subset],axis=1)
full_subset.set_index('SalePrice')
full_subset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2919 entries, 0 to 2918
Data columns (total 49 columns):
SalePrice      1460 non-null float64
OverallQual    2919 non-null int64
```

GrLivArea	2919	non-null	float64
GarageCars	2918	non-null	float64
GarageArea	2918	non-null	float64
1stFlrSF	2919	non-null	float64
FullBath	2919	non-null	int64
YearBuilt	2919	non-null	int64
YearRemodAdd	2919	non-null	int64
GarageYrBltd	2760	non-null	float64
TotRmsAbvGrd	2919	non-null	float64
Fireplaces	2919	non-null	int64
OpenPorchSF	2919	non-null	float64
MasVnrArea	2896	non-null	float64
LotArea	2919	non-null	float64
TotalBsmtSF	2918	non-null	float64
LotFrontage	2433	non-null	float64
WoodDeckSF	2919	non-null	float64
HalfBath	2919	non-null	int64
BsmtFullBath	2917	non-null	float64
BedroomAbvGr	2919	non-null	int64
BsmtUnfSF	2918	non-null	float64
BsmtFinSF1	2918	non-null	float64
2ndFlrSF	2919	non-null	float64
ScreenPorch	2919	non-null	float64
Neighborhood_E	2919	non-null	float64
ExterQual_E	2919	non-null	float64
KitchenQual_E	2919	non-null	float64
BsmtQual_E	2919	non-null	float64
GarageFinish_E	2919	non-null	float64
GarageType_E	2919	non-null	float64
Foundation_E	2919	non-null	float64
FireplaceQu_E	2919	non-null	float64
HeatingQC_E	2919	non-null	float64
MasVnrType_E	2919	non-null	float64
Exterior1st_E	2919	non-null	float64
MSZoning_E	2919	non-null	float64
Exterior2nd_E	2919	non-null	float64
BsmtFinType1_E	2919	non-null	float64
GarageCond_E	2919	non-null	float64
GarageQual_E	2919	non-null	float64
BsmtExposure_E	2919	non-null	float64
CentralAir_E	2919	non-null	float64
SaleType_E	2919	non-null	float64
HouseStyle_E	2919	non-null	float64
Electrical_E	2919	non-null	float64
PavedDrive_E	2919	non-null	float64
BsmtCond_E	2919	non-null	float64
LotShape_E	2919	non-null	float64

dtypes: float64(42), int64(7)

memory usage: 1.1 MB

6h-1. Train Subset

```
In [147]: train_subset = full_subset[:1460]
          train_subset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 49 columns):
SalePrice          1460 non-null float64
OverallQual        1460 non-null int64
GrLivArea          1460 non-null float64
GarageCars         1460 non-null float64
GarageArea         1460 non-null float64
1stFlrSF           1460 non-null float64
FullBath           1460 non-null int64
YearBuilt          1460 non-null int64
YearRemodAdd       1460 non-null int64
GarageYrBlt        1379 non-null float64
TotRmsAbvGrd      1460 non-null float64
Fireplaces         1460 non-null int64
OpenPorchSF        1460 non-null float64
MasVnrArea         1452 non-null float64
LotArea           1460 non-null float64
TotalBsmtSF        1460 non-null float64
LotFrontage        1201 non-null float64
WoodDeckSF         1460 non-null float64
HalfBath           1460 non-null int64
BsmtFullBath       1460 non-null float64
BedroomAbvGr       1460 non-null int64
BsmtUnfSF          1460 non-null float64
BsmtFinSF1         1460 non-null float64
2ndFlrSF           1460 non-null float64
ScreenPorch        1460 non-null float64
Neighborhood_E     1460 non-null float64
ExterQual_E        1460 non-null float64
KitchenQual_E      1460 non-null float64
BsmtQual_E         1460 non-null float64
GarageFinish_E     1460 non-null float64
GarageType_E       1460 non-null float64
Foundation_E       1460 non-null float64
FireplaceQu_E      1460 non-null float64
HeatingQC_E        1460 non-null float64
MasVnrType_E       1460 non-null float64
Exterior1st_E      1460 non-null float64
MSZoning_E         1460 non-null float64
```



```

Exterior2nd_E      1460 non-null float64
BsmtFinType1_E     1460 non-null float64
GarageCond_E       1460 non-null float64
GarageQual_E       1460 non-null float64
BsmtExposure_E     1460 non-null float64
CentralAir_E       1460 non-null float64
SaleType_E         1460 non-null float64
HouseStyle_E       1460 non-null float64
Electrical_E       1460 non-null float64
PavedDrive_E       1460 non-null float64
BsmtCond_E         1460 non-null float64
LotShape_E         1460 non-null float64
dtypes: float64(42), int64(7)
memory usage: 559.0 KB

```

6h-3. Test Subset

```

In [148]: test_subset = full_subset[1460: ]
          test_subset.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1459 entries, 1460 to 2918
Data columns (total 49 columns):
SalePrice          0 non-null float64
OverallQual        1459 non-null int64
GrLivArea          1459 non-null float64
GarageCars         1458 non-null float64
GarageArea         1458 non-null float64
1stFlrSF           1459 non-null float64
FullBath           1459 non-null int64
YearBuilt          1459 non-null int64
YearRemodAdd       1459 non-null int64
GarageYrBlt        1381 non-null float64
TotRmsAbvGrd       1459 non-null float64
Fireplaces         1459 non-null int64
OpenPorchSF        1459 non-null float64
MasVnrArea         1444 non-null float64
LotArea            1459 non-null float64
TotalBsmtSF        1458 non-null float64
LotFrontage        1232 non-null float64
WoodDeckSF         1459 non-null float64
HalfBath           1459 non-null int64
BsmtFullBath       1457 non-null float64
BedroomAbvGr       1459 non-null int64
BsmtUnfSF          1458 non-null float64
BsmtFinSF1         1458 non-null float64
2ndFlrSF           1459 non-null float64

```

```

ScreenPorch      1459 non-null float64
Neighborhood_E   1459 non-null float64
ExterQual_E      1459 non-null float64
KitchenQual_E    1459 non-null float64
BsmtQual_E       1459 non-null float64
GarageFinish_E   1459 non-null float64
GarageType_E     1459 non-null float64
Foundation_E     1459 non-null float64
FireplaceQu_E    1459 non-null float64
HeatingQC_E      1459 non-null float64
MasVnrType_E     1459 non-null float64
Exterior1st_E    1459 non-null float64
MSZoning_E       1459 non-null float64
Exterior2nd_E    1459 non-null float64
BsmtFinType1_E   1459 non-null float64
GarageCond_E     1459 non-null float64
GarageQual_E     1459 non-null float64
BsmtExposure_E   1459 non-null float64
CentralAir_E     1459 non-null float64
SaleType_E       1459 non-null float64
HouseStyle_E     1459 non-null float64
Electrical_E     1459 non-null float64
PavedDrive_E     1459 non-null float64
BsmtCond_E       1459 non-null float64
LotShape_E       1459 non-null float64
dtypes: float64(42), int64(7)
memory usage: 558.6 KB

```

6h-1. Viewing Features for Predictive Model

```
In [149]: train_subset.head(20)
```

```

Out[149]:
   SalePrice  OverallQual  GrLivArea  GarageCars  GarageArea  1stFlrSF  \
0   12.247699           7    7.444833           2.0         548.0    6.753438
1   12.109016           6    7.141245           2.0         460.0    7.141245
2   12.317171           7    7.488294           2.0         608.0    6.825460
3   11.849405           7    7.448916           3.0         642.0    6.869014
4   12.429220           8    7.695758           3.0         836.0    7.044033
5   11.870607           5    7.217443           2.0         480.0    6.680855
6   12.634606           8    7.435438           2.0         636.0    7.435438
7   12.206078           7    7.645398           2.0         484.0    7.010312
8   11.774528           7    7.481556           2.0         468.0    6.930495
9   11.678448           5    6.982863           1.0         205.0    6.982863
10  11.771444           5    6.947937           1.0         384.0    6.947937
11  12.751303           9    7.751475           3.0         736.0    7.075809
12  11.877576           5    6.816736           1.0         352.0    6.816736
13  12.540761           7    7.309881           3.0         840.0    7.309881

```

14	11.964007	6	7.134094	1.0	352.0	7.134094
15	11.790565	7	6.751101	2.0	576.0	6.751101
16	11.911708	6	6.912743	2.0	480.0	6.912743
17	11.407576	4	7.167809	2.0	516.0	7.167809
18	11.976666	5	7.016610	2.0	576.0	7.016610
19	11.842236	5	7.200425	1.0	294.0	7.200425

	FullBath	YearBuilt	YearRemodAdd	GarageYrBlt	...	GarageCond_E \
0	2	2003	2003	2003.0	...	6.0
1	2	1976	1976	1976.0	...	6.0
2	2	2001	2002	2001.0	...	6.0
3	1	1915	1970	1998.0	...	6.0
4	2	2000	2000	2000.0	...	6.0
5	1	1993	1995	1993.0	...	6.0
6	2	2004	2005	2004.0	...	6.0
7	2	1973	1973	1973.0	...	6.0
8	2	1931	1950	1931.0	...	6.0
9	1	1939	1950	1939.0	...	6.0
10	1	1965	1965	1965.0	...	6.0
11	3	2005	2006	2005.0	...	6.0
12	1	1962	1962	1962.0	...	6.0
13	2	2006	2007	2006.0	...	6.0
14	1	1960	1960	1960.0	...	6.0
15	1	1929	2001	1991.0	...	6.0
16	1	1970	1970	1970.0	...	6.0
17	2	1967	1967	1967.0	...	6.0
18	1	2004	2004	2004.0	...	6.0
19	1	1958	1965	1958.0	...	6.0

	GarageQual_E	BsmtExposure_E	CentralAir_E	SaleType_E	HouseStyle_E \
0	4.0	2.0	2.0	5.0	7.0
1	4.0	5.0	2.0	5.0	5.0
2	4.0	3.0	2.0	5.0	7.0
3	4.0	2.0	2.0	5.0	7.0
4	4.0	4.0	2.0	5.0	7.0
5	4.0	2.0	2.0	5.0	3.0
6	4.0	4.0	2.0	5.0	5.0
7	4.0	3.0	2.0	5.0	7.0
8	3.0	2.0	2.0	5.0	3.0
9	6.0	2.0	2.0	5.0	1.0
10	4.0	2.0	2.0	5.0	5.0
11	4.0	2.0	2.0	8.0	7.0
12	4.0	2.0	2.0	5.0	5.0
13	4.0	4.0	2.0	8.0	5.0
14	4.0	2.0	2.0	5.0	5.0
15	4.0	2.0	2.0	5.0	1.0
16	4.0	2.0	2.0	5.0	5.0
17	4.0	1.0	2.0	5.0	5.0

18	4.0	2.0	2.0	5.0	5.0
19	4.0	2.0	2.0	4.0	5.0

	Electrical_E	PavedDrive_E	BsmtCond_E	LotShape_E
0	6.0	3.0	4.0	1.0
1	6.0	3.0	4.0	1.0
2	6.0	3.0	4.0	2.0
3	6.0	3.0	5.0	2.0
4	6.0	3.0	4.0	2.0
5	6.0	3.0	4.0	2.0
6	6.0	3.0	4.0	1.0
7	6.0	3.0	4.0	2.0
8	3.0	3.0	4.0	1.0
9	6.0	3.0	4.0	1.0
10	6.0	3.0	4.0	1.0
11	6.0	3.0	4.0	2.0
12	6.0	3.0	4.0	4.0
13	6.0	3.0	4.0	2.0
14	6.0	3.0	4.0	2.0
15	4.0	3.0	4.0	1.0
16	6.0	3.0	4.0	2.0
17	6.0	3.0	2.0	1.0
18	6.0	3.0	4.0	1.0
19	6.0	3.0	4.0	1.0

[20 rows x 49 columns]

6h-2. Inspecting New DataFrame

```
In [150]: train_subset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 49 columns):
SalePrice      1460 non-null float64
OverallQual    1460 non-null int64
GrLivArea      1460 non-null float64
GarageCars     1460 non-null float64
GarageArea     1460 non-null float64
1stFlrSF       1460 non-null float64
FullBath       1460 non-null int64
YearBuilt      1460 non-null int64
YearRemodAdd   1460 non-null int64
GarageYrBltd   1379 non-null float64
TotRmsAbvGrd   1460 non-null float64
Fireplaces     1460 non-null int64
OpenPorchSF    1460 non-null float64
MasVnrArea     1452 non-null float64
```

LotArea	1460	non-null	float64
TotalBsmtSF	1460	non-null	float64
LotFrontage	1201	non-null	float64
WoodDeckSF	1460	non-null	float64
HalfBath	1460	non-null	int64
BsmtFullBath	1460	non-null	float64
BedroomAbvGr	1460	non-null	int64
BsmtUnfSF	1460	non-null	float64
BsmtFinSF1	1460	non-null	float64
2ndFlrSF	1460	non-null	float64
ScreenPorch	1460	non-null	float64
Neighborhood_E	1460	non-null	float64
ExterQual_E	1460	non-null	float64
KitchenQual_E	1460	non-null	float64
BsmtQual_E	1460	non-null	float64
GarageFinish_E	1460	non-null	float64
GarageType_E	1460	non-null	float64
Foundation_E	1460	non-null	float64
FireplaceQu_E	1460	non-null	float64
HeatingQC_E	1460	non-null	float64
MasVnrType_E	1460	non-null	float64
Exterior1st_E	1460	non-null	float64
MSZoning_E	1460	non-null	float64
Exterior2nd_E	1460	non-null	float64
BsmtFinType1_E	1460	non-null	float64
GarageCond_E	1460	non-null	float64
GarageQual_E	1460	non-null	float64
BsmtExposure_E	1460	non-null	float64
CentralAir_E	1460	non-null	float64
SaleType_E	1460	non-null	float64
HouseStyle_E	1460	non-null	float64
Electrical_E	1460	non-null	float64
PavedDrive_E	1460	non-null	float64
BsmtCond_E	1460	non-null	float64
LotShape_E	1460	non-null	float64

dtypes: float64(42), int64(7)
memory usage: 559.0 KB

1.6.7 7. Feature Engineering

7a. has_feature Function

```
In [151]: def has_feature(feature):
           if feature > 0:
               return 1
           else:
               return 0
```

7b. Full Data Set Engineering

```
In [152]: full_subset['HasBasement'] = full_subset['TotalBsmtSF'].apply(has_feature)
full_subset['HasGarage'] = full_subset['GarageArea'].apply(has_feature)
full_subset['Has2ndFlr'] = full_subset['2ndFlrSF'].apply(has_feature)
full_subset['HasOpenPorch'] = full_subset['OpenPorchSF'].apply(has_feature)
full_subset['HasScreenPorch'] = full_subset['ScreenPorch'].apply(has_feature)
full_subset['HasFirePlace'] = full_subset['Fireplaces'].apply(has_feature)
full_subset['HasMsVnr'] = full_subset['MasVnrArea'].apply(has_feature)
full_subset['HasWoodDeck'] = full_subset['WoodDeckSF'].apply(has_feature)
full_subset['HasPool'] = houses_full['PoolArea'].apply(has_feature)
full_subset['HasBasementBathroom'] = full_subset['BsmtFullBath'].apply(has_feature)
full_subset['ExtraBathrooms'] = full_subset['FullBath'].apply(lambda x: 1 if x > 1 else 0)
full_subset['New'] = full_subset['YearBuilt'].apply(lambda x: 1 if x > 2000 else 0)
full_subset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2919 entries, 0 to 2918
```

```
Data columns (total 61 columns):
```

SalePrice	1460	non-null	float64
OverallQual	2919	non-null	int64
GrLivArea	2919	non-null	float64
GarageCars	2918	non-null	float64
GarageArea	2918	non-null	float64
1stFlrSF	2919	non-null	float64
FullBath	2919	non-null	int64
YearBuilt	2919	non-null	int64
YearRemodAdd	2919	non-null	int64
GarageYrBltd	2760	non-null	float64
TotRmsAbvGrd	2919	non-null	float64
Fireplaces	2919	non-null	int64
OpenPorchSF	2919	non-null	float64
MasVnrArea	2896	non-null	float64
LotArea	2919	non-null	float64
TotalBsmtSF	2918	non-null	float64
LotFrontage	2433	non-null	float64
WoodDeckSF	2919	non-null	float64
HalfBath	2919	non-null	int64
BsmtFullBath	2917	non-null	float64
BedroomAbvGr	2919	non-null	int64
BsmtUnfSF	2918	non-null	float64
BsmtFinSF1	2918	non-null	float64
2ndFlrSF	2919	non-null	float64
ScreenPorch	2919	non-null	float64
Neighborhood_E	2919	non-null	float64
ExterQual_E	2919	non-null	float64
KitchenQual_E	2919	non-null	float64
BsmtQual_E	2919	non-null	float64
GarageFinish_E	2919	non-null	float64

```

GarageType_E          2919 non-null float64
Foundation_E          2919 non-null float64
FireplaceQu_E         2919 non-null float64
HeatingQC_E           2919 non-null float64
MasVnrType_E          2919 non-null float64
Exterior1st_E         2919 non-null float64
MSZoning_E            2919 non-null float64
Exterior2nd_E         2919 non-null float64
BsmtFinType1_E        2919 non-null float64
GarageCond_E          2919 non-null float64
GarageQual_E          2919 non-null float64
BsmtExposure_E        2919 non-null float64
CentralAir_E          2919 non-null float64
SaleType_E            2919 non-null float64
HouseStyle_E          2919 non-null float64
Electrical_E          2919 non-null float64
PavedDrive_E          2919 non-null float64
BsmtCond_E            2919 non-null float64
LotShape_E            2919 non-null float64
HasBasement           2919 non-null int64
HasGarage              2919 non-null int64
Has2ndFlr             2919 non-null int64
HasOpenPorch          2919 non-null int64
HasScreenPorch        2919 non-null int64
HasFirePlace          2919 non-null int64
HasMsVnr              2919 non-null int64
HasWoodDeck           2919 non-null int64
HasPool               2919 non-null int64
HasBasementBathroom   2919 non-null int64
ExtraBathrooms        2919 non-null int64
New                   2919 non-null int64
dtypes: float64(42), int64(19)
memory usage: 1.4 MB

```

7c. Train_subset Engineering

```

In [153]: train_subset = full_subset[:1460]
          train_subset.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 61 columns):
SalePrice          1460 non-null float64
OverallQual        1460 non-null int64
GrLivArea          1460 non-null float64
GarageCars         1460 non-null float64
GarageArea         1460 non-null float64

```

1stFlrSF	1460	non-null	float64
FullBath	1460	non-null	int64
YearBuilt	1460	non-null	int64
YearRemodAdd	1460	non-null	int64
GarageYrBltn	1379	non-null	float64
TotRmsAbvGrd	1460	non-null	float64
Fireplaces	1460	non-null	int64
OpenPorchSF	1460	non-null	float64
MasVnrArea	1452	non-null	float64
LotArea	1460	non-null	float64
TotalBsmtSF	1460	non-null	float64
LotFrontage	1201	non-null	float64
WoodDeckSF	1460	non-null	float64
HalfBath	1460	non-null	int64
BsmtFullBath	1460	non-null	float64
BedroomAbvGr	1460	non-null	int64
BsmtUnfSF	1460	non-null	float64
BsmtFinSF1	1460	non-null	float64
2ndFlrSF	1460	non-null	float64
ScreenPorch	1460	non-null	float64
Neighborhood_E	1460	non-null	float64
ExterQual_E	1460	non-null	float64
KitchenQual_E	1460	non-null	float64
BsmtQual_E	1460	non-null	float64
GarageFinish_E	1460	non-null	float64
GarageType_E	1460	non-null	float64
Foundation_E	1460	non-null	float64
FireplaceQu_E	1460	non-null	float64
HeatingQC_E	1460	non-null	float64
MasVnrType_E	1460	non-null	float64
Exterior1st_E	1460	non-null	float64
MSZoning_E	1460	non-null	float64
Exterior2nd_E	1460	non-null	float64
BsmtFinType1_E	1460	non-null	float64
GarageCond_E	1460	non-null	float64
GarageQual_E	1460	non-null	float64
BsmtExposure_E	1460	non-null	float64
CentralAir_E	1460	non-null	float64
SaleType_E	1460	non-null	float64
HouseStyle_E	1460	non-null	float64
Electrical_E	1460	non-null	float64
PavedDrive_E	1460	non-null	float64
BsmtCond_E	1460	non-null	float64
LotShape_E	1460	non-null	float64
HasBasement	1460	non-null	int64
HasGarage	1460	non-null	int64
Has2ndFlr	1460	non-null	int64
HasOpenPorch	1460	non-null	int64


```

HasScreenPorch      1460 non-null int64
HasFirePlace        1460 non-null int64
HasMsVnr            1460 non-null int64
HasWoodDeck         1460 non-null int64
HasPool             1460 non-null int64
HasBasementBathroom 1460 non-null int64
ExtraBathrooms      1460 non-null int64
New                 1460 non-null int64
dtypes: float64(42), int64(19)
memory usage: 695.9 KB

```

7d. Test Set

```

In [154]: test_subset = full_subset[1460: ]
          test_subset.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1459 entries, 1460 to 2918
Data columns (total 61 columns):
SalePrice      0 non-null float64
OverallQual    1459 non-null int64
GrLivArea      1459 non-null float64
GarageCars     1458 non-null float64
GarageArea     1458 non-null float64
1stFlrSF       1459 non-null float64
FullBath       1459 non-null int64
YearBuilt      1459 non-null int64
YearRemodAdd   1459 non-null int64
GarageYrBltd   1381 non-null float64
TotRmsAbvGrd   1459 non-null float64
Fireplaces     1459 non-null int64
OpenPorchSF    1459 non-null float64
MasVnrArea     1444 non-null float64
LotArea        1459 non-null float64
TotalBsmtSF    1458 non-null float64
LotFrontage    1232 non-null float64
WoodDeckSF     1459 non-null float64
HalfBath       1459 non-null int64
BsmtFullBath   1457 non-null float64
BedroomAbvGr   1459 non-null int64
BsmtUnfSF      1458 non-null float64
BsmtFinSF1     1458 non-null float64
2ndFlrSF       1459 non-null float64
ScreenPorch    1459 non-null float64
Neighborhood_E 1459 non-null float64
ExterQual_E    1459 non-null float64
KitchenQual_E  1459 non-null float64

```



BsmtQual_E	1459	non-null	float64
GarageFinish_E	1459	non-null	float64
GarageType_E	1459	non-null	float64
Foundation_E	1459	non-null	float64
FireplaceQu_E	1459	non-null	float64
HeatingQC_E	1459	non-null	float64
MasVnrType_E	1459	non-null	float64
Exterior1st_E	1459	non-null	float64
MSZoning_E	1459	non-null	float64
Exterior2nd_E	1459	non-null	float64
BsmtFinType1_E	1459	non-null	float64
GarageCond_E	1459	non-null	float64
GarageQual_E	1459	non-null	float64
BsmtExposure_E	1459	non-null	float64
CentralAir_E	1459	non-null	float64
SaleType_E	1459	non-null	float64
HouseStyle_E	1459	non-null	float64
Electrical_E	1459	non-null	float64
PavedDrive_E	1459	non-null	float64
BsmtCond_E	1459	non-null	float64
LotShape_E	1459	non-null	float64
HasBasement	1459	non-null	int64
HasGarage	1459	non-null	int64
Has2ndFlr	1459	non-null	int64
HasOpenPorch	1459	non-null	int64
HasScreenPorch	1459	non-null	int64
HasFirePlace	1459	non-null	int64
HasMsVnr	1459	non-null	int64
HasWoodDeck	1459	non-null	int64
HasPool	1459	non-null	int64
HasBasementBathroom	1459	non-null	int64
ExtraBathrooms	1459	non-null	int64
New	1459	non-null	int64


dtypes: float64(42), int64(19)
memory usage: 695.4 KB



1.6.8 8. Visualizing Data



8a. Scatter Plots of Selected Features



```
In [155]: for feature in train_subset:
            plt.scatter(train_subset[feature], train_subset['SalePrice'])
            plt.title('SalePrice vs ' + feature)
            plt.ylabel('SalePrice')
            plt.xlabel(feature)
            plt.show()
```

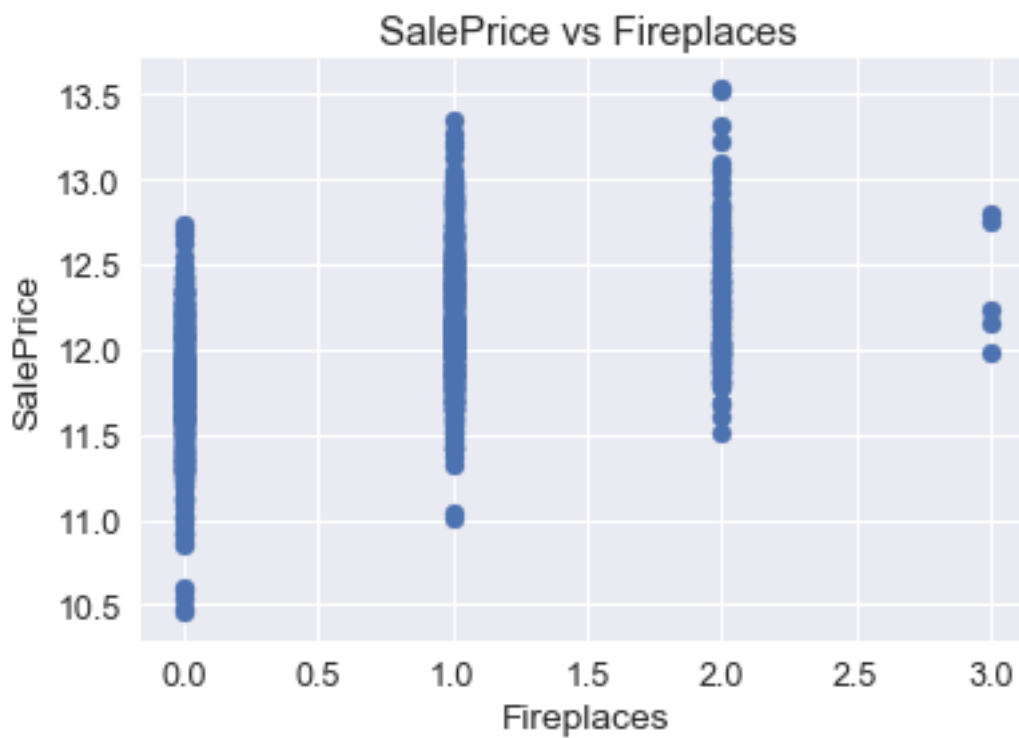












































































































8b. Removing Outliers from GrLivArea

```
In [156]: train_subset = train_subset.drop(train_subset[train_subset['GrLivArea'] > 8.3].index
          train_subset[train_subset['SalePrice'] < 12.25].index)

train_subset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1458 entries, 0 to 1459
Data columns (total 61 columns):
SalePrice           1458 non-null float64
OverallQual         1458 non-null int64
GrLivArea           1458 non-null float64
GarageCars          1458 non-null float64
GarageArea          1458 non-null float64
1stFlrSF            1458 non-null float64
FullBath            1458 non-null int64
YearBuilt           1458 non-null int64
YearRemodAdd        1458 non-null int64
GarageYrBlt         1377 non-null float64
TotRmsAbvGrd        1458 non-null float64
Fireplaces          1458 non-null int64
OpenPorchSF         1458 non-null float64
MasVnrArea          1450 non-null float64
```


LotArea	1458	non-null	float64
TotalBsmtSF	1458	non-null	float64
LotFrontage	1199	non-null	float64
WoodDeckSF	1458	non-null	float64
HalfBath	1458	non-null	int64
BsmtFullBath	1458	non-null	float64
BedroomAbvGr	1458	non-null	int64
BsmtUnfSF	1458	non-null	float64
BsmtFinSF1	1458	non-null	float64
2ndFlrSF	1458	non-null	float64
ScreenPorch	1458	non-null	float64
Neighborhood_E	1458	non-null	float64
ExterQual_E	1458	non-null	float64
KitchenQual_E	1458	non-null	float64
BsmtQual_E	1458	non-null	float64
GarageFinish_E	1458	non-null	float64
GarageType_E	1458	non-null	float64
Foundation_E	1458	non-null	float64
FireplaceQu_E	1458	non-null	float64
HeatingQC_E	1458	non-null	float64
MasVnrType_E	1458	non-null	float64
Exterior1st_E	1458	non-null	float64
MSZoning_E	1458	non-null	float64
Exterior2nd_E	1458	non-null	float64
BsmtFinType1_E	1458	non-null	float64
GarageCond_E	1458	non-null	float64
GarageQual_E	1458	non-null	float64
BsmtExposure_E	1458	non-null	float64
CentralAir_E	1458	non-null	float64
SaleType_E	1458	non-null	float64
HouseStyle_E	1458	non-null	float64
Electrical_E	1458	non-null	float64
PavedDrive_E	1458	non-null	float64
BsmtCond_E	1458	non-null	float64
LotShape_E	1458	non-null	float64
HasBasement	1458	non-null	int64
HasGarage	1458	non-null	int64
Has2ndFlr	1458	non-null	int64
HasOpenPorch	1458	non-null	int64
HasScreenPorch	1458	non-null	int64
HasFirePlace	1458	non-null	int64
HasMsVnr	1458	non-null	int64
HasWoodDeck	1458	non-null	int64
HasPool	1458	non-null	int64
HasBasementBathroom	1458	non-null	int64
ExtraBathrooms	1458	non-null	int64
New	1458	non-null	int64

dtypes: float64(42), int64(19)

memory usage: 706.2 KB

8b-1. SalePrice vs GrLivArea Scatterplot with Outliers Removed

```
In [157]: plt.scatter(train_subset['GrLivArea'], train_subset['SalePrice'])
plt.title('SalePrice vs. GrLivArea')
plt.ylabel('SalePrice')
plt.xlabel('GrLivArea')
plt.show()
```



1.6.9 9. Getting Predictions

9a. Getting Dummy Variables

```
In [158]: houses_full = pd.get_dummies(houses_full)
train_subset = pd.get_dummies(train_subset)
```

9b. Filling Missing Values

```
In [159]: houses_full = houses_full.fillna(houses_full.mean())
train_subset = train_subset.fillna(houses_full.mean())
test_subset = test_subset.fillna(houses_full.mean())
```

9c. Partitioning Data

```
In [160]: X_train, X_test, y_train, y_test = train_test_split(train_subset.drop('SalePrice', axis=1), train_subset['SalePrice'],
                                                            test_size = 0.3, random_state = 3)
```

9d. Writing Function to Record RMSE

```
In [161]: # Define error measure for official scoring : RMSE
scorer = make_scorer(mean_squared_error, greater_is_better = False)

def rmse_cv_train(model):
    rmse= np.sqrt(-cross_val_score(model, X_train, y_train, scoring = scorer, cv = 10))
    return(rmse)

def rmse_cv_test(model):
    rmse= np.sqrt(-cross_val_score(model, X_test, y_test, scoring = scorer, cv = 10))
    return(rmse)
```

9e. Linear Regression Model


```
In [162]: lr = LinearRegression()
lr.fit(X_train, y_train)


# Look at predictions on training and validation set
print("RMSE on Training set :", rmse_cv_train(lr).mean())
print("RMSE on Test set :", rmse_cv_test(lr).mean())
y_train_pred = lr.predict(X_train)
y_test_pred = lr.predict(X_test)

# Plot residuals
plt.scatter(y_train_pred, y_train_pred - y_train, c = "blue", marker = "s", label = "Training residuals")
plt.scatter(y_test_pred, y_test_pred - y_test, c = "lightgreen", marker = "s", label = "Validation residuals")
plt.title("Linear regression")
plt.xlabel("Predicted values")
plt.ylabel("Residuals")
plt.legend(loc = "upper left")
plt.hlines(y = 0, xmin = 10.5, xmax = 13.5, color = "red")
plt.show()

# Plot predictions
plt.scatter(y_train_pred, y_train, c = "blue", marker = "s", label = "Training data")
plt.scatter(y_test_pred, y_test, c = "lightgreen", marker = "s", label = "Validation data")
plt.title("Linear regression")
plt.xlabel("Predicted values")
plt.ylabel("Real values")
plt.legend(loc = "upper left")
plt.plot([10.5, 13.5], [10.5, 13.5], c = "red")
plt.show()
```

RMSE on Training set : 0.124804458115
RMSE on Test set : 0.133658974553





9f. Ridge Model

In [163]: *# 2* Ridge*

```
ridge = RidgeCV(alphas = [0.01, 0.03, 0.06, 0.1, 0.3, 0.6, 1, 3, 6, 10, 30, 60])
ridge.fit(X_train, y_train)
alpha = ridge.alpha_
print("Best alpha :", alpha)

print("Try again for more precision with alphas centered around " + str(alpha))
ridge = RidgeCV(alphas = [alpha * .6, alpha * .65, alpha * .7, alpha * .75, alpha *
                        alpha * .9, alpha * .95, alpha, alpha * 1.05, alpha * 1.1,
                        alpha * 1.25, alpha * 1.3, alpha * 1.35, alpha * 1.4],
                cv = 10)
ridge.fit(X_train, y_train)
alpha = ridge.alpha_
print("Best alpha :", alpha)

print("Ridge RMSE on Training set :", rmse_cv_train(ridge).mean())
print("Ridge RMSE on Test set :", rmse_cv_test(ridge).mean())
y_train_rdg = ridge.predict(X_train)
y_test_rdg = ridge.predict(X_test)

# Plot residuals
```

```

plt.scatter(y_train_rdg, y_train_rdg - y_train, c = "blue", marker = "s", label = "T")
plt.scatter(y_test_rdg, y_test_rdg - y_test, c = "lightgreen", marker = "s", label = "V")
plt.title("Linear regression with Ridge regularization")
plt.xlabel("Predicted values")
plt.ylabel("Residuals")
plt.legend(loc = "upper left")
plt.hlines(y = 0, xmin = 10.5, xmax = 13.5, color = "red")
plt.show()

# Plot predictions
plt.scatter(y_train_rdg, y_train, c = "blue", marker = "s", label = "Training data")
plt.scatter(y_test_rdg, y_test, c = "lightgreen", marker = "s", label = "Validation data")
plt.title("Linear regression with Ridge regularization")
plt.xlabel("Predicted values")
plt.ylabel("Real values")
plt.legend(loc = "upper left")
plt.plot([10.5, 13.5], [10.5, 13.5], c = "red")
plt.show()

# Plot important coefficients
coefs = pd.Series(ridge.coef_, index = X_train.columns)
print("Ridge picked " + str(sum(coefs != 0)) + " features and eliminated the other " +
      str(sum(coefs == 0)) + " features")
imp_coefs = pd.concat([coefs.sort_values().head(10),
                      coefs.sort_values().tail(10)])
imp_coefs.plot(kind = "barh")
plt.title("Coefficients in the Ridge Model")
plt.show()

```



Best alpha : 0.3

Try again for more precision with alphas centered around 0.3


Best alpha : 0.195

Ridge RMSE on Training set : 0.124772422841

Ridge RMSE on Test set : 0.131839691941



Ridge picked 60 features and eliminated the other 0 features



9g. Lasso Model

In [164]: # 3* Lasso

```
lasso = LassoCV(alphas = [0.0001, 0.0003, 0.0006, 0.001, 0.003, 0.006, 0.01, 0.03, 0.06, 0.1, 0.3, 0.6, 1],
                  max_iter = 50000, cv = 10)
lasso.fit(X_train, y_train)
alpha = lasso.alpha_
print("Best alpha :", alpha)

print("Try again for more precision with alphas centered around " + str(alpha))
lasso = LassoCV(alphas = [alpha * .6, alpha * .65, alpha * .7, alpha * .75, alpha * .8,
                           alpha * .85, alpha * .9, alpha * .95, alpha, alpha * 1.05,
                           alpha * 1.1, alpha * 1.15, alpha * 1.25, alpha * 1.3, alpha * 1.4],
                  max_iter = 50000, cv = 10)
lasso.fit(X_train, y_train)
alpha = lasso.alpha_
print("Best alpha :", alpha)

print("Lasso RMSE on Training set :", rmse_cv_train(lasso).mean())
print("Lasso RMSE on Test set :", rmse_cv_test(lasso).mean())
y_train_las = lasso.predict(X_train)
y_test_las = lasso.predict(X_test)
```

```

# Plot residuals
plt.scatter(y_train_las, y_train_las - y_train, c = "blue", marker = "s", label = "Training data")
plt.scatter(y_test_las, y_test_las - y_test, c = "lightgreen", marker = "s", label = "Validation data")
plt.title("Linear regression with Lasso regularization")
plt.xlabel("Predicted values")
plt.ylabel("Residuals")
plt.legend(loc = "upper left")
plt.hlines(y = 0, xmin = 10.5, xmax = 13.5, color = "red")
plt.show()

# Plot predictions
plt.scatter(y_train_las, y_train, c = "blue", marker = "s", label = "Training data")
plt.scatter(y_test_las, y_test, c = "lightgreen", marker = "s", label = "Validation data")
plt.title("Linear regression with Lasso regularization")
plt.xlabel("Predicted values")
plt.ylabel("Real values")
plt.legend(loc = "upper left")
plt.plot([10.5, 13.5], [10.5, 13.5], c = "red")
plt.show()

# Plot important coefficients
coefs = pd.Series(lasso.coef_, index = X_train.columns)
print("Lasso picked " + str(sum(coefs != 0)) + " features and eliminated the other " +
      str(sum(coefs == 0)) + " features")
imp_coefs = pd.concat([coefs.sort_values().head(10),
                      coefs.sort_values().tail(10)])
imp_coefs.plot(kind = "barh")
plt.title("Coefficients in the Lasso Model")
plt.show()

```



Best alpha : 0.0001

Try again for more precision with alphas centered around 0.0001


Best alpha : 6e-05

Lasso RMSE on Training set : 0.124920699681

Lasso RMSE on Test set : 0.132494703915



Lasso picked 59 features and eliminated the other 1 features



9h. ElasticNet Model

```
In [165]: elasticNet = ElasticNetCV(l1_ratio = [0.1, 0.3, 0.5, 0.6, 0.7, 0.8, 0.85, 0.9, 0.95,
                                         alphas = [0.0001, 0.0003, 0.0006, 0.001, 0.003, 0.006,
                                                    0.01, 0.03, 0.06, 0.1, 0.3, 0.6, 1, 3, 6],
                                         max_iter = 50000, cv = 10)

elasticNet.fit(X_train, y_train)
alpha = elasticNet.alpha_
ratio = elasticNet.l1_ratio_
print("Best l1_ratio :", ratio)
print("Best alpha :", alpha )

print("Try again for more precision with l1_ratio centered around " + str(ratio))
elasticNet = ElasticNetCV(l1_ratio = [ratio * .85, ratio * .9, ratio * .95, ratio, ratio * 1.05],
                           alphas = [0.0001, 0.0003, 0.0006, 0.001, 0.003, 0.006, 0.01, 0.03, 0.06, 0.1, 0.3, 0.6, 1, 3, 6],
                           max_iter = 50000, cv = 10)

elasticNet.fit(X_train, y_train)
if (elasticNet.l1_ratio_ > 1):
    elasticNet.l1_ratio_ = 1
alpha = elasticNet.alpha_
ratio = elasticNet.l1_ratio_
print("Best l1_ratio :", ratio)
print("Best alpha :", alpha )
```

```

print("Now try again for more precision on alpha, with l1_ratio fixed at " + str(ratio)
      " and alpha centered around " + str(alpha))
elasticNet = ElasticNetCV(l1_ratio = ratio,
                          alphas = [alpha * .6, alpha * .65, alpha * .7, alpha * .75,
                                     alpha * .95, alpha, alpha * 1.05, alpha * 1.1, alpha * 1.2,
                                     alpha * 1.35, alpha * 1.4],
                          max_iter = 50000, cv = 10)
elasticNet.fit(X_train, y_train)
if (elasticNet.l1_ratio_ > 1):
    elasticNet.l1_ratio_ = 1
alpha = elasticNet.alpha_
ratio = elasticNet.l1_ratio_
print("Best l1_ratio :", ratio)
print("Best alpha :", alpha )

print("ElasticNet RMSE on Training set :", rmse_cv_train(elasticNet).mean())
print("ElasticNet RMSE on Test set :", rmse_cv_test(elasticNet).mean())
y_train_ela = elasticNet.predict(X_train)
y_test_ela = elasticNet.predict(X_test)


# Plot residuals
plt.scatter(y_train_ela, y_train_ela - y_train, c = "blue", marker = "s", label = "Training residuals")
plt.scatter(y_test_ela, y_test_ela - y_test, c = "lightgreen", marker = "s", label = "Validation residuals")
plt.title("Linear regression with ElasticNet regularization")
plt.xlabel("Predicted values")
plt.ylabel("Residuals")
plt.legend(loc = "upper left")
plt.hlines(y = 0, xmin = 10.5, xmax = 13.5, color = "red")
plt.show()


# Plot predictions
plt.scatter(y_train, y_train_ela, c = "blue", marker = "s", label = "Training data")
plt.scatter(y_test, y_test_ela, c = "lightgreen", marker = "s", label = "Validation data")
plt.title("Linear regression with ElasticNet regularization")
plt.xlabel("Predicted values")
plt.ylabel("Real values")
plt.legend(loc = "upper left")
plt.plot([10.5, 13.5], [10.5, 13.5], c = "red")
plt.show()

# Plot important coefficients
coefs = pd.Series(elasticNet.coef_, index = X_train.columns)
print("ElasticNet picked " + str(sum(coefs != 0)) + " features and eliminated the others")
imp_coefs = pd.concat([coefs.sort_values().head(10),
                      coefs.sort_values().tail(10)])
imp_coefs.plot(kind = "barh")
plt.title("Coefficients in the ElasticNet Model")
plt.show()


```

Best l1_ratio : 0.1
Best alpha : 0.0001
Try again for more precision with l1_ratio centered around 0.1
Best l1_ratio : 0.085
Best alpha : 0.0003
Now try again for more precision on alpha, with l1_ratio fixed at 0.085 and alpha centered around 0.0001
Best l1_ratio : 0.085
Best alpha : 0.000195
ElasticNet RMSE on Training set : 0.124827915458
ElasticNet RMSE on Test set : 0.132185935611





ElasticNet picked 59 features and eliminated the other 1 features



1.6.10 10. Kaggle Predictions

10a. Dropping Target Variable

```
In [166]: test_subset = test_subset.drop(['SalePrice'], axis = 1)
```

10b. Getting Predictions

```
In [167]: predictions = lasso.predict(test_subset)
          predictions = np.expm1(predictions)
```

10c. Creating Submission DataFrame and CSV

```
In [168]: ID = test['Id']
          submission = pd.DataFrame({'Id' : ID, 'SalePrice': predictions})
          submission.to_csv('SalePrice_Predictions.csv', index=False)
          submission.head()
```

```
Out[168]:
```

	Id	SalePrice
1460	1461	123014.788596
1461	1462	159440.482579
1462	1463	182106.821506
1463	1464	193958.762850
1464	1465	207943.672650